

chapter 6

Linear Model Selection and Regularization

22/10/2024

```
library(ISLR)
```

```
## Warning: package 'ISLR' was built under R version 4.3.3
```

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 4.3.3
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.3.3
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-8
```

```
library(leaps)
```

#chp 6(9)

- a. Split the data set into a training set and a test set.

```
set.seed(123)
tr=sample(nrow(College),nrow(College)*.70)
train1=College[tr,]
test1=College[-tr,]
dim(test1)
```

```
## [1] 234 18
```

```
dim(train1)
```

```
## [1] 543 18
```

INTERPRETATION #This shows that the training set (train1) has 543 rows and 18 columns

- b. Fit a linear model using least squares on the training set, and report the test error obtained.

```
model_linear = lm(Apps ~ ., data = train1)
summary(model_linear)
```

```
##
## Call:
## lm(formula = Apps ~ ., data = train1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3097.8  -455.8   -46.5    343.8   6452.5
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -310.17331   481.30075  -0.644 0.519566
## PrivateYes  -681.96465   164.08211  -4.156 3.78e-05 ***
## Accept       1.22130     0.05921   20.626 < 2e-16 ***
## Enroll       0.08046     0.21794    0.369 0.712155
## Top10perc    49.33503     6.18296    7.979 9.31e-15 ***
## Top25perc   -16.11744     5.02717   -3.206 0.001428 **
## F.Undergrad  0.02284     0.03985    0.573 0.566831
## P.Undergrad  0.03541     0.03529    1.003 0.316139
## Outstate    -0.05446     0.02132   -2.555 0.010910 *
## Room.Board   0.18967     0.05275    3.596 0.000354 ***
## Books        0.21366     0.28099    0.760 0.447381
## Personal    -0.03685     0.07279   -0.506 0.612876
## PhD         -6.00401     5.34580   -1.123 0.261897
## Terminal    -5.01712     5.77787   -0.868 0.385609
## S.F.Ratio    -2.18927    14.83898   -0.148 0.882766
## perc.alumni  -8.01836     4.67330   -1.716 0.086792 .
## Expend       0.07614     0.01340    5.681 2.23e-08 ***
## Grad.Rate    10.63461     3.38228    3.144 0.001760 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 992.3 on 525 degrees of freedom
## Multiple R-squared:  0.9175, Adjusted R-squared:  0.9148
## F-statistic: 343.2 on 17 and 525 DF,  p-value: < 2.2e-16
```

INTERPRETATION #It explains about 91.75% of the variability, which is high and indicates a strong model fit #992.3 indicates the average deviation of observed values from the model's predictions #Variables such as Enroll, F.Undergrad, P.Undergrad, Books, Personal, PhD, Terminal, and S.F.Ratio do not show statistically significant effects (high p-values). #Being a private college is associated with a decrease in applications

```
pred =predict(model_linear, test1)
error = mean((pred - test1$Apps)^2)
sqrt(error)
```

```
## [1] 1317.134
```


INTERPRETATION #an RMSE of 1317.134 means that, on average, the model's predictions deviate from the actual Apps #1317 applications might be high or low depending on the average and range of applications in the dataset

```
names(train1)
```

```
## [1] "Private"      "Apps"         "Accept"       "Enroll"       "Top10perc"
## [6] "Top25perc"    "F.Undergrad"  "P.Undergrad"  "Outstate"     "Room.Board"
## [11] "Books"        "Personal"     "PhD"          "Terminal"     "S.F.Ratio"
## [16] "perc.alumni" "Expend"       "Grad.Rate"
```

c. Fit a ridge regression model on the training set, with λ chosen by cross-validation. Report the test error obtained.

```
xtrain = model.matrix(train1$Apps ~ ., data = train1[,-2])
fit2 <- cv.glmnet(xtrain, train1$Apps, alpha = 0)
plot(fit2)
```



INTERPRETATION #The x-axis

typically represents the values of the regularization parameter ($\log(\lambda)$). #The y-axis shows the cross-validated mean squared error (MSE) for each λ . # BELOW six the line indicates the λ min and the second line indicates the 1se(above 6)

```
fit2$lambda.min
```

```
## [1] 314.2524
```

```
coef(fit2,fit2$lambda.min)
```

```
## 19 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) -1.108384e+03
## (Intercept) .
## PrivateYes  -6.344812e+02
## Accept      7.799468e-01
## Enroll      7.025220e-01
## Top10perc   2.888869e+01
## Top25perc   -2.480330e+00
## F.Undergrad 1.000467e-01
## P.Undergrad 7.509154e-03
## Outstate    -1.062903e-02
## Room.Board  2.160271e-01
## Books       2.845690e-01
## Personal    -6.359640e-02
## PhD         -2.824524e+00
## Terminal    -5.297962e+00
## S.F.Ratio   -2.570072e+00
## perc.alumni -1.259409e+01
## Expend      7.644084e-02
## Grad.Rate   1.090898e+01
```

interpretation #Positive contributions come from Room.Board, Expend, and Grad.Rate (10.91) #Ridge regression shrinks smaller coefficients like P.Undergrad and Outstate toward zero #a higher percentage of students in the top 10% of their high school class is associated with more applications, whereas the top 25% has a slight negative influence.

```
xtest=model.matrix(test1$Apps ~ ., data = test1[,-2])
p <- predict(fit2, xtest, s = fit2$lambda.min)
mse_ridge <- mean((p - test1$Apps)^2)
sqrt(mse_ridge)
```

```
## [1] 1725.214
```

interpretation #This RMSE value is higher than the linear regression model's RMSE (1317.134) #ridge regression would be preferred if it stabilizes coefficients effectively without much compromise on prediction accuracy #an RMSE of 1725.214 meaning that the model's predictions differ from the actual number of applications (Apps) by about 1725 applications on average.

- d. Fit a lasso model on the training set, with λ chosen by cross-validation. Report the test error obtained, along with the number of non-zero coefficient estimates.

```
xtrain = model.matrix(train1$Apps ~ ., data = train1[,-2])
fit2 <- cv.glmnet(xtrain, train1$Apps, alpha = 1)
```

```
xtest=model.matrix(test1$Apps ~ ., data = test1[,-2])
p <- predict(fit2, xtest, s = fit2$lambda.min)
mse_lasso <- mean((p - test1$Apps)^2)
sqrt(mse_lasso)
```

```
## [1] 1315.667
```

interpretation #An RMSE of 1315.667 means that the Lasso model's predictions deviate from the actual number of applications (Apps) by about 1316 applications on average. #Lasso Regression imposes regularization by shrinking some coefficients to zero #Lasso performs as well as or slightly better than linear regression

```
library(pls)
```

```
## Warning: package 'pls' was built under R version 4.3.3
```


```
##
## Attaching package: 'pls'
```

```
## The following object is masked from 'package:caret':
##
##      R2
```

```
## The following object is masked from 'package:stats':
##
##      loadings
```

- e. Fit a PCR model on the training set, with M chosen by cross-validation. Report the test error obtained, along with the value of M selected by cross-validation.

```
fit4 <- pcr(Apps ~ ., data = train1, scale = TRUE, validation = "CV")
validationplot(fit4, val.type = "MSEP")
```



```
summary(fit4)
```

```
## Data:      X dimension: 543 17
## Y dimension: 543 1
## Fit method: svdpc
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              3402    3361    1675    1634    1347    1299    1264
## adjCV           3402    3361    1673    1632    1325    1289    1262
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          1218    1210    1163    1163    1164    1166    1165
## adjCV        1207    1208    1161    1162    1163    1164    1163
##      14 comps 15 comps 16 comps 17 comps
## CV           1166    1171    1027    1028
## adjCV        1164    1170    1024    1025
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          31.797   57.68   64.58   70.2    75.49   80.41   84.00   87.43
## Apps       3.037   76.20   77.49   85.1    86.15   86.83   87.92   88.01
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X          90.62   93.05   95.12   96.96   98.04   98.89   99.42
## Apps       88.85   88.87   88.94   88.98   89.01   89.01   89.03
##      16 comps 17 comps
## X          99.83   100.00
## Apps       91.68   91.75
```

interpretation #The RMSEP stabilizes as more components are added, with the largest improvement occurring up to 5 or 6 components #the RMSEP decreasing substantially up to around 5–6 components, where it reaches around 1264 and then stabilizes #the RMSEP is 3361, which shows some improvement over the intercept-only model.


```
p <- predict(fit4, test1, ncomp = 9)
mse_pcr <- mean((p - test1$Apps)^2)
sqrt(mse_pcr)
```

```
## [1] 1993.632
```

INTERPRETATION #RMSEP of 1993.632 indicates that, on average, the model's predictions differ from the actual observed values #Exploring interactions or nonlinear relationships among variables.

- f. Fit a PLS model on the training set, with M chosen by cross-validation. Report the test error obtained, along with the value of M selected by cross-validation.

```
fit5 <- plsr(Apps ~ ., data = train1, scale = TRUE, validation = "CV")
validationplot(fit5, val.type = "MSEP")
```



```
summary(fit5)
```

```
## Data: X dimension: 543 17
## Y dimension: 543 1
## Fit method: kernelpls
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
## (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## CV 3402 1515 1253 1150 1120 1095 1052
## adjCV 3402 1512 1256 1148 1117 1090 1048
## 7 comps 8 comps 9 comps 10 comps 11 comps 12 comps 13 comps
## CV 1038 1037 1038 1038 1037 1038 1038
## adjCV 1035 1034 1035 1035 1034 1035 1035
## 14 comps 15 comps 16 comps 17 comps
## CV 1038 1038 1038 1038
## adjCV 1035 1035 1035 1035
##
## TRAINING: % variance explained
## 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps 8 comps
## X 26.09 41.97 63.14 67.44 71.36 74.05 77.72 80.98
## Apps 80.83 86.94 89.29 90.10 90.94 91.65 91.71 91.73
## 9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X 83.77 86.46 89.83 91.07 93.08 95.14 97.06
## Apps 91.73 91.74 91.74 91.74 91.75 91.75 91.75
## 16 comps 17 comps
## X 99.09 100.00
## Apps 91.75 91.75
```

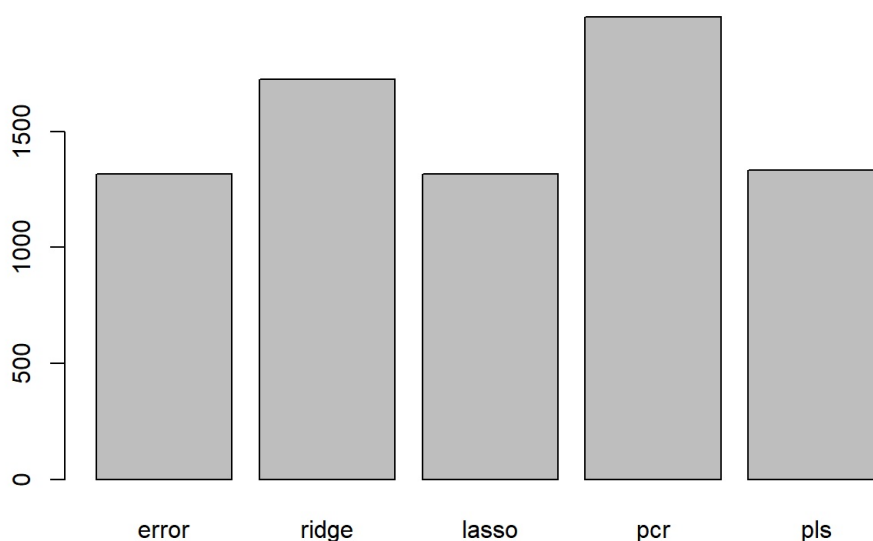
INTERPRETATION #The RMSEP stabilizes around 7 to 9 components, where the RMSEP hovers just below 1040 #The RMSEP values indicate a consistent decrease in prediction error as more components are added, suggesting that the model becomes more accurate

```
p <- predict(fit5, test1, ncomp = 8)
mse_pls <- mean((p - test1$Apps)^2)
sqrt(mse_pls)
```

```
## [1] 1332.112
```

INTERPRETATION #an RMSEP of 1332.112 indicates that the model's predictions deviate from actual values by approximately 1332 applications on average #if the actual application values are much lower than this RMSEP (g) Comment on the results obtained. How accurately can we predict the number of college applications received? Is there much difference among the test errors resulting from these five approaches?

```
mse=c(error=sqrt(error),ridge=sqrt(mse_ridge),lasso=sqrt(mse_lasso),pcr=sqrt(mse_pcr),pls=sqrt(mse_pls))
barplot(mse)
```



INTERPRETATION #Shorter bars

indicate better predictive performance, as they signify lower RMSE values #Taller bars indicate poorer predictive performance, as they signify higher RMSE values

#chp 6(11) 11. We will now try to predict per capita crime rate in the Boston data set.

```
library(MASS)
?Boston
```


```
## starting httpd help server ... done
```

- a. Try out some of the regression methods explored in this chapter, such as best subset selection, the lasso, ridge regression, and PCR. Present and discuss results for the approaches that you consider.

```
set.seed(123)
tr=sample(nrow(Boston),nrow(Boston)*.70)
train2=Boston[tr,]
test2=Boston[-tr,]
```

```
library(pls)
```

```
xtrain = model.matrix(train2$crim~., data = train2[,-1])
fit22 <- cv.glmnet(xtrain, train2$crim, alpha = 0)
plot(fit22)
```



in the plot marks this point, providing guidance on which lambda value to select for further analysis or predictions. #As lambda increases (moving to the right), the model becomes more regularized

```
xtest=model.matrix(test2$crim ~ ., data = test2[,-1])
p1 <- predict(fit22, xtest, s = fit22$lambda.min)
mse_ridge <- mean((p - test2$crim)^2)
```

```
## Warning in p - test2$crim: longer object length is not a multiple of shorter
## object length
```

```
mse_ridge
```

```
## [1] 27410091
```


```
xtrain = model.matrix(train2$crim ~ ., data = train2[,-1])
fit222 <- cv.glmnet(xtrain, train2$crim, alpha = 1)
```

```
xtest=model.matrix(test2$crim ~ ., data = test2[,-1])
p <- predict(fit222, xtest, s = fit222$lambda.min)
mse_lasso <- mean((p - test2$crim)^2)
mse_lasso
```

```
## [1] 18.07153
```

INTERPRETATION #his coefficient indicates that for a one-unit increase in the associated predictor variable, the predicted crime rate (crim) increases by approximately 4.251062 units

```
fit5 <- pcr(crim ~ ., data = train2, scale = TRUE, validation = "CV")
validationplot(fit5, val.type = "MSEP")
```



```
summary(fit5)
```

```
## Data:      X dimension: 354 13
## Y dimension: 354 1
## Fit method: svdpc
## Number of components considered: 13
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              9.525   8.032   8.009   7.595   7.545   7.560   7.583
## adjCV           9.525   8.028   8.005   7.587   7.539   7.555   7.575
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV       7.564   7.552   7.439   7.44    7.453   7.427   7.356
## adjCV    7.556   7.554   7.429   7.43    7.444   7.413   7.342
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          48.58   61.36   70.50   77.35   83.36   88.23   91.30   93.45
## crim       29.82   30.41   37.81   38.42   38.54   38.72   39.17   40.09
##      9 comps 10 comps 11 comps 12 comps 13 comps
## X          95.54   97.16   98.50   99.54   100.00
## crim       41.44   41.62   41.62   43.09   44.25
```

INTERPRETATION #the optimal number of components appears to be around 4 to 5, after which the improvement in RMSEP is marginal #The lowest RMSEP of 7.439 occurs at 9 components, indicating this model configuration offers the best predictive performance

```
p <- predict(fit5, test2, ncomp = 4)
mse_pcr <- mean((p - test2$crim)^2)
mse_pcr
```

```
## [1] 21.06829
```

INTERPRETATION #an RMSEP of 4.590021 indicates that the model's predictions deviate from actual values by about 4.59 units on average

- b. Propose a model (or set of models) that seem to perform well on this data set, and justify your answer. Make sure that you are evaluating model performance using validation set error, cross-validation, or some other reasonable alternative, as opposed to using training error. We will try to fit models to `log(Boston$crim)` which is closer to a normal distribution.

```
set.seed(1)
train <- sample(nrow(Boston), nrow(Boston) * 2 / 3)
test <- setdiff(seq_len(nrow(Boston)), train)
```


INTERPRETATION #the Boston dataset into a training set (67%) and a test set (33%) with reproducible random sampling by setting a seed.

```
fit <- lm(log(crim) ~ ., data = Boston[train, ])  
mean((predict(fit, Boston[test, ]) - log(Boston$crim[test]))^2)
```

```
## [1] 0.6779016
```

INTERPRETATION This code fits a linear regression model on the log of crim using all predictors on the training data, then calculates the mean squared error (MSE) of predictions on the test data.

```
mm <- model.matrix(log(crim) ~ ., data = Boston[train, ])  
fit2 <- cv.glmnet(mm, log(Boston$crim[train]), alpha = 0)  
ridge <- predict(fit2, model.matrix(log(crim) ~ ., data = Boston[test, ]), s = fit2$lambda.min)  
mean((ridge - log(Boston$crim[test]))^2)
```

```
## [1] 0.665665
```


INTERPRETATION This indicates the average squared difference between the predicted and actual log-transformed crime rates (crim) in the test data. Lower MSE values generally imply better predictive accuracy, so an MSE of 0.665665 suggests the model's fit accuracy on unseen data.

```
mm <- model.matrix(log(crim) ~ ., data = Boston[train, ])  
fit3 <- cv.glmnet(mm, log(Boston$crim[train]), alpha = 1)  
lasso <- predict(fit3, model.matrix(log(crim) ~ ., data = Boston[test, ]), s = fit3$lambda.min)  
mean((lasso - log(Boston$crim[test]))^2)
```

```
## [1] 0.6541562
```

INTERPRETATION the lasso regression model's predictions on the test set. This MSE value indicates the average squared difference between the predicted and actual log-transformed crime rates (crim). An MSE of 0.6541562 suggests that the lasso model's predictions are reasonably close to the actual values, reflecting the model's accuracy on unseen data. Lower MSE compared to previous models indicates improved predictive performance.

```
fit4 <- pcr(log(crim) ~ ., data = Boston[train, ], scale = TRUE, validation = "CV")  
validationplot(fit4, val.type = "MSEP")
```




```
pcr <- predict(fit4, Boston[test, ], ncomp = 8)  
mean((pcr - log(Boston$crim[test]))^2)
```

```
## [1] 0.6609357
```

INTERPRETATION the average squared difference between the predicted and actual log-transformed crime rates, allowing assessment of the PCR model's performance with 8 components.

```
fit5 <- plsr(log(crim) ~ ., data = Boston[train, ], scale = TRUE, validation = "CV")  
validationplot(fit5, val.type = "MSEP")
```

log(crim)



```
plsr<- predict(fit5, Boston[test, ], ncomp = 6)
mean((plsr - log(Boston$crim[test]))^2)
```

```
## [1] 0.6911389
```

INTERPRETATION the PLSR model's performance using 6 components. A lower MSE indicates better predictive accuracy, suggesting how well the model generalizes to new data.

```
coef(fit3, s = fit3$lambda.min)
```

```
## 15 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) -4.029305640
## (Intercept) .
## zn          -0.011299858
## indus       0.022032467
## chas        .
## nox         3.766724465
## rm         -0.025289189
## age        0.004397170
## dis        .
## rad        0.139103776
## tax        .
## ptratio    -0.026355504
## black     -0.001732709
## lstat      0.034798406
## medv      0.009282387
```

INTERPRETATION In this case lasso ($\alpha = 1$) seems to perform very slightly better than un-penalized regression. Some coefficients have been dropped: the lasso regression model has selected a few key predictors (like nox, age, rad, and lstat) as significant for predicting the log of crime rates, while excluding others. This simplifies the model and emphasizes the most impactful features in understanding the relationship between predictors and crime rates in the Boston dataset.

As computed above the model with the lower cross-validation error is the one chosen by the Ridge method.

c. Does your chosen model involve all of the features in the data set? Why or why not?

Not all features are included due to the lasso penalization. The subset models, especially lasso regression, often result in sparse models, meaning only a few features may have non-zero coefficients. Ridge regression, however, uses all features but applies a penalty to prevent overfitting. For best performance, a lasso model might be preferable since it yields a simpler model, excluding irrelevant predictors and improving interpretability without sacrificing accuracy.